

Memory Architectures for the Agentic AI Era:

Extreme Hardware-Software Co-Design for Heterogeneous AI Factories

Author: Ping Long, Chief Systems Architect, Lead Researcher, [SiliconLanguage](#)

Contact: [LinkedIn](#) | [GitHub](#) | [0kernel.ai](#) | plongpingl@gmail.com | fellow@0kernel.ai

Abstract and Executive Summary

The transition toward agentic Artificial Intelligence (AI), multimodal reasoning pipelines, and trillion-parameter Large Language Models (LLMs) has fundamentally exposed the physical and architectural limitations of traditional Von Neumann computing paradigms. Over the past decade, semiconductor scaling has allowed computational units—particularly the Tensor Cores and Warpgroup Matrix-Multiply-Accumulate (WGMMMA) engines within modern Graphics Processing Units (GPUs)—to scale their arithmetic logic capabilities into the realm of hundreds of teraflops per die. However, the surrounding systems architecture has encountered a severe and restrictive "Memory Wall." In contemporary AI factory deployments, the performance bottleneck has shifted entirely away from arithmetic logic unit (ALU) saturation and toward the complex logistics of data movement. System throughput is now dictated almost exclusively by memory bandwidth, non-deterministic access latency, and the physical interconnect topologies binding disparate compute nodes together. ^[1]

To address these compounding constraints, the rigid historical dichotomy between byte-addressable volatile memory (managed by the CPU hardware) and block-addressable persistent storage (managed by the operating system kernel) has definitively dissolved. Modern hyperscale infrastructure now relies on a continuous, multi-tier fabric that abstracts geographical location, physical interconnect boundaries, and media persistence. This paradigm shift requires a fundamental hardware-software co-design approach, where zero-kernel data planes directly manipulate the physical memory controllers of remote nodes over massive Ethernet and InfiniBand fabrics.

This comprehensive research report rigorously analyzes the contemporary memory and storage taxonomy, validating the evolving boundaries of Compute Express Link (CXL) 3.1 and 4.0 attached memory against traditional storage-class tiers, and formally introducing Tier 2.5 (G2.5) Persistent Context Memory. Furthermore, it details the state-of-the-art in storage and network bypassing, focusing on the Data Plane Development Kit (DPDK), the Storage Performance Development Kit (SPDK), user-space memory mapping, universal AI compilation, and the precise orchestration of the Key-Value (KV) cache staging hierarchy across distributed inference systems. ^[2] By examining cutting-edge physical specifications—including JEDEC's recently ratified High Bandwidth Memory 4 (HBM4) and Low-Power Double Data Rate 6 (LPDDR6) standards, alongside algorithmic breakthroughs, open interconnects like UALink, and heterogeneous compute vectors—this document provides a highly detailed, actionable blueprint for the engineering of advanced, memory-disaggregated AI factories.

1. The Baseline Physical Memory Taxonomy, Disaggregated

Fabrics, and Zero-Kernel Data Planes

The baseline taxonomy categorizing memory architectures into a rigid continuum must be continuously updated to reflect the ratification of advanced specifications, specifically the CXL 3.1 standard, and the deployment of advanced fabric-attached topologies. The architectural separation between Tier 2 (Fabric-Attached Volatile Memory) and Tier 3 (Storage-Class Persistent Memory) has been heavily scrutinized by systems engineers.

The validated, linear tiering system, integrating current-generation metrics and the new G2.5 tier, is structured as follows:

Fabric-Attached Memory and Zero-Kernel Data Planes

To address this multi-tier fabric, operating systems are transitioning to zero-kernel data planes. Frameworks like the Fabric-Attached Memory File System (FAMFS) operate as scale-out warm memory file systems over 100 Gb/s fabrics, exposing CXL 3.1/4.0 Dynamic Capacity Devices (DCDs) to user-space without incurring VFS/POSIX overheads.^[3] FAMFS maps CXL device allocations directly to user-space VMAs via Direct Access (DAX), enabling cache-line level access without requiring end-user application modifications. Metadata changes are coordinated via an append-only log written by a Master node and "played" by Client nodes, eliminating per-host write-back races.^[3]

By combining PCIe peer-to-peer DMA with the Data Plane Development Kit (DPDK) and the Storage Performance Development Kit (SPDK), applications utilize lock-free Shared Work Queues (SWQs) to poll for memory completions, bypassing the OS kernel entirely to process over 100 million IOPS. SPDK achieves extreme CPU efficiency, saturating hardware arrays using just a single core under native benchmarking tools, whereas asynchronous kernel frameworks like `io_uring` require multiple threads and incur high CPU overhead.

To maintain strict multi-tenant security within these disaggregated, bare-metal planes, silicon-edge isolation is required. Modern SSD controller DMA engines natively enforce this by parsing hardware-level Process Address Space IDs (PASID) and executing Logical Block Address (LBA) Range Access Control directly on the media endpoint, ensuring secure isolation without invoking the host CPU.

Architectural Implications of CXL 3.1 and 4.0 Integration

The CXL 3.1 specification radically alters multi-host resource disaggregation and hardware-enforced coherency. Earlier iterations of the CXL standard provided rudimentary point-to-point capacity expansion over PCIe physical layers, largely limited to single-node hierarchy-based routing. CXL 3.1 transitions this into a fully scale-out, multi-level switched fabric architecture capable of supporting non-tree topologies, including mesh, ring, star, and multi-dimensional butterfly networks.^[4]

To maintain absolute coherency as memory becomes pooled and shared across multiple heterogeneous computational nodes, CXL 3.1 enforces hardware-level back-invalidate semantics. When a peer device updates a specific 64-byte cache line residing inside a shared CXL.mem target, the memory controller automatically issues upstream back-invalidate messages to all hosts and accelerators that have previously cached that line.^[5]

Furthermore, Host Physical Address Device-Managed Coherence (HDM-DB) delegates the complex tracking of cache line states directly to the memory device rather than relying on the host OS or the host CPU's home agent.^[5] Accelerators are now granted the capability to cache host memory directly, supported by up to 32 bits of metadata injected into each 64-byte cache line.^[5]

The Next Frontier: CXL 4.0 and PCIe 7.0 Integration Building on the advancements of CXL 3.1, the CXL 4.0

specification (ratified in November 2025) further revolutionizes the fabric by integrating the PCIe 7.0 physical layer. By doubling the Nyquist frequency while preserving the PAM4 signaling and 256-byte FLIT-based encoding of PCIe 6.0, PCIe 7.0 and CXL 4.0 deliver a raw data rate of 128 GT/s, achieving up to 512 GB/s of bidirectional bandwidth per x16 link without incurring any additional protocol latency. The most critical architectural innovation in CXL 4.0 is the introduction of "Bundled Ports." This feature allows infrastructure engineers to aggregate multiple physical CXL links into a single logical connection, enabling link bandwidths of up to 1.5 TB/s. This extreme throughput is the foundational enabler for true multi-rack memory pooling and massive-scale disaggregation across the data center.

Eliminating the OS Kernel: SPDK, DPDK, and io_uring

Modern AI architectures rely entirely on zero-kernel data planes to interact from Tier1 to Tier 4, enabling applications to achieve end-to-end zero-copy data routing.

The Storage Performance Development Kit (SPDK) eliminates interrupt latency by utilizing lock-free, user-space polling mechanisms.^[6] Instead of relying on the Linux kernel to schedule I/O and sleep the thread until a hardware interrupt signals NVMe completion, the application spins a dedicated thread pinned to a specific CPU core. This thread continuously polls the PCIe memory-mapped registers of the NVMe drive. While this consumes 100% of the assigned CPU core's cycles, it guarantees deterministic, microsecond-level tail latencies and massive IOPS throughput by avoiding context switching and software lock contention entirely.^[6] In parallel, the Data Plane Development Kit (DPDK) bypasses the traditional TCP/IP network stack, allowing network packets to be placed directly into application buffers or GPU VRAM (via GPUDirect RDMA) without the host CPU ever touching the payload.^[7]

Hardware Acceleration and Innovation

Navigating the memory wall also requires extreme bare-metal optimizations tailored to the specific microarchitectures of NVIDIA and AMD accelerators.

Extreme Co-Design: Heterogeneous Architecture in the Rubin-LPX Era

The integration of the **NVIDIA Vera Rubin** platform with the **Groq 3 LPX** rack represents a shift in NVIDIA's hardware strategy toward "extreme co-design" and heterogeneous architecture.^[44] By treating the entire NVL72 rack as a single AI supercomputer, NVIDIA optimizes every phase of the inference loop to eliminate latency, maximize throughput, and reduce energy consumption.^[44, 45]

Here is a summary of how the Vera Rubin + LPX pipeline functions as a series of targeted hardware optimizations:

Attention-FFN Disaggregation (AFD) for Workload Splitting

Instead of forcing a single GPU to handle the entire inference workload, execution is physically split token-by-token.^[45, 46] **Rubin GPUs** use their compute density to handle the context-heavy prefill phase and attention mechanisms, while **Groq 3 LPUs** are dedicated to latency-sensitive Feed-Forward Network (FFN) and Mixture of Experts (MoE) layers.^[45]

Fusion Memory Architecture

As AI models process increasingly long contexts, the platform leverages two completely different memory paradigms simultaneously to avoid bottlenecks:^[46]

- **Dynamic KV Cache in HBM4:** The expanding Key-Value (KV) cache is kept in the Rubin GPU's massive

288GB of HBM4, which delivers 22 TB/s of bandwidth (as established in our previous conversation).^[47]

- **Static Weights in SRAM:** FFN weight matrices are pinned permanently inside the Groq 3 LPX rack's 128GB of flat, cacheless SRAM, which allows computations to operate without off-chip memory retrieval delays.^[45, 46]

Deterministic Execution and Speculative Decoding

The Groq 3 LPU is optimized for strict, compiler-orchestrated determinism, eliminating runtime jitter and providing perfectly stable token generation.^[46] NVIDIA leverages this extreme speed for speculative decoding: the LPX rack acts as a high-speed draft-generation engine to guess tokens rapidly, while the powerful Rubin GPUs run parallel verification to finalize the tokens.^[45]

Highly Specialized Interconnect Fabrics

To ensure the split architecture does not bottleneck at the network level, NVIDIA utilizes distinct, highly optimized data paths:

- **NVLink 6:** Within the Rubin rack, GPUs communicate over a bidirectional all-to-all fabric, allowing massive MoE models to share data across 72 GPUs as if they were a single accelerator.
- **NVLink-C2C:** The Rubin GPUs are paired with the custom **Vera CPU** (an ARM Olympus processor) to offload data staging and agentic orchestration so the GPUs are never starved for data.^[45, 47]

The Ultimate Result

By physically disaggregating the inference loop and optimizing the memory and interconnects, the Vera Rubin + LPX platform achieves up to 35x higher inference throughput per megawatt and drives a 10x reduction in the cost per token compared to the previous Blackwell generation (metrics detailed in our earlier discussion).^[44]

AMD MoRI-IO KV Cache Connector

Concurrently, for AMD Instinct environments (such as the MI355X), the ecosystem has introduced the Modular RDMA Interface (MoRI), specifically the MoRI-IO KV cache connector for vLLM. By leveraging specialized "communication bubbles," MoRI-IO overlaps KV transfer with active computation. During the prefill phase, as each transformer layer completes its calculation, the MoRI-IO connector pushes the KV blocks directly into the VRAM of the decode instance via RDMA.^[8] This zero-copy pipeline essentially removes the prefill communication bottleneck, maximizing total distributed throughput for massive Mixture-of-Experts (MoE) topologies.^[8]

2. The Agentic Inference Context Framework, KV Cache Disaggregation, and Algorithmic Breakthroughs

The Inference Context Memory Hierarchy

The necessity of strict memory tiering, hardware-enforced coherency, and zero-copy data paths culminates in the execution of agentic Large Language Models. Unlike static inference tasks, agentic, multi-turn AI requires maintaining a mathematical representation of all prior conversational context and attention states—known as the Key-Value (KV) cache. The physical capacity required to store the KV cache scales linearly with sequence length and batch size, modeled algorithmically by the formula:

$$\text{Capacity}_{\text{KV}} = 2 \times L \times N_{\text{heads}} \times D_{\text{head}} \times S_{\text{seq}} \times \text{Precision}$$

From a rapid infrastructure planning perspective, systems engineers utilize heuristic approximations to calculate the context footprint based on input sequence length (ISL):

$$\text{Context Size} \approx \text{ISL} \times \text{Memory per Token} \times \text{Quantization Factor}$$

At 128K context lengths with concurrent batching, the KV cache alone can rapidly demand upwards of 160 GB, pushing it far beyond the capacity bounds of a single Tier 0 HBM package.

Because the GPU's High Bandwidth Memory (HBM, Tier 0) is prohibitively expensive and strictly capacity-limited, inference clusters must aggressively page the KV cache out to lower memory tiers ^[12].

Algorithmic Memory Breakthroughs: Subquadratic Sparse Attention (SSA)

Scaling the context window to millions of tokens traditionally demands an exponential increase in VRAM capacity, primarily due to the mathematical complexity of standard dense attention mechanisms. To combat this, the industry introduced Subquadratic Sparse Attention (SSA) architectures in mid-2026. By implementing content-dependent sparse attention, SSA completely bypasses the all-pairs quadratic scaling, achieving a linear growth curve for the KV cache. This yields extraordinary efficiency gains, such as a 62.5x reduction in attention FLOPs at a 1-million token sequence compared to baseline quadratic models. Consequently, models utilizing SSA can natively reason over ultra-long context windows of up to 12 million tokens without exhausting the 100GB+ VRAM bounds of modern Tier 0 and Tier 1 hardware arrays.

Managing the G1 through G4 Storage Continuum

To sustain this scale, the KV cache must be paged across a rigid storage continuum, explicitly layering out to include high-speed pod-level extensions:

- **Tier 1 (G1) - Active KV (GPU HBM):** The hottest context, reserved strictly for tokens actively attended to by the compute cores.
- **Tier 2 (G2) - Staging KV (System DRAM):** Temporary eviction hold for immediately reusable blocks.
- **Tier 2.5 (G2.5) - Persistent Context Memory:** Utilizing CXL Type-3 devices, this byte-addressable persistent staging layer acts as a high-speed overflow for KV cache fragments before they suffer the latency penalty of hitting traditional block-I/O endpoints.
- **Tier 3 (G3) - Warm KV (Node-Local NVMe):** Local SSD arrays providing high-capacity, microsecond-latency block storage.
- **Tier 3.5 (G3.5) - Fabric Context Extension:** An Ethernet-attached, pod-level flash tier optimized for hot, reusable KV cache across the entire compute pod.
- **Tier 4 (G4) - Cold / Shared KV (Networked Storage):** The ground-truth layer for long-term state persistence.

The Architectural Evolution: Tier G3.5 and Context Memory eXtension (CMX)

NVIDIA formally defined **Tier G3.5** through the **Context Memory eXtension (CMX)** architecture.^[9] CMX operates as an Ethernet-attached, pod-level flash tier. Powered by the NVIDIA BlueField-4 Data Processing Unit and Spectrum-X Ethernet switches, CMX virtualizes dense NVMe SSD arrays into a massive, globally accessible Key-Value storage layer.^[10] The host GPU nodes treat CMX as an extended local cache using the **DOCA Memos API**. The BlueField-4 DPU acts as the "storage brain," offloading routing and hardware-accelerated encryption, allowing the compute cores to focus entirely on token generation.^[9]

The Next Architectural Evolution: Tier G2.5 and Elastic CXL Persistent Memory

To bridge the performance and volatility gap between host DRAM and node-local NVMe, this report defines **Tier G2.5** as the domain of **CXL Type-3 persistent memory devices**.^[54] This tier integrates high-density persistent media with a cache-coherent, byte-addressable host interface via the CXL.mem protocol. Unlike traditional block-storage tiers, Tier G2.5 functions as a native extension of the memory subsystem, enabling a highly durable, byte-addressable staging ground that fundamentally transforms the data plane for agentic inference.

1. Hardware-Software Co-Design for Wear-Free Logging

The primary architectural utility of Tier G2.5 is the enablement of **flash-wear-free transaction logging**. By mapping database Write-Ahead Logs (WALs) and high-frequency KV cache metadata directly into the G2.5 byte-addressable persistent address space, the system utilizes native CPU load/store semantics. This approach completely bypasses block-based I/O overhead—such as *fsync* system calls—and eliminates the destructive write amplification typically associated with routing fragmented, random KV updates to traditional NVMe NAND flash.

2. Protocol Integration and Data Plane Determinism

Tier G2.5 devices minimize tail latency by bypassing the legacy NVMe stack entirely. Leveraging the CXL.mem protocol, the runtime maps these devices into the coherent address space, allowing for:

- **Byte-Addressable Durability:** The system employs explicit CPU memory barriers (e.g., *clwb* and *sfence*) to commit log entries with nanosecond-level determinism, ensuring data consistency without triggering the heavy overhead of OS block-layer drivers.
- **Deterministic Reasoning Loops via Flash-Bypass:** By migrating KV cache management to Tier G2.5's byte-addressable persistent memory, the architecture eliminates the stochastic latency "jitter" typically introduced by SSD controller background tasks—such as garbage collection and wear-leveling cycles. This transition from block-I/O to memory-semantic access not only guarantees the deterministic throughput required for long-context agentic reasoning loops but also prevents the repetitive NAND flash wear-out that occurs when high-frequency token updates are routed to traditional storage.

3. Heterogeneous Staging Strategy

Tier G2.5 operates as a critical staging orchestrator in the heterogeneous memory hierarchy. It acts as an **augmentation tier**, not a replacement for local High Bandwidth Memory (HBM). While HBM provides the extreme bandwidth necessary for active tensor computation, Tier G2.5 provides the **elastic memory capacity** required to hold massive KV cache states that exceed local HBM limits. It serves as a transient buffer for short-lived, block-aligned KV cache fragments before they are coalesced by the runtime into large, sequential chunks and flushed to Tier G3 (Bulk NVMe). This tier-aware strategy ensures that only bulk model weights and archival context reach the high-density flash tier, effectively isolating the NVMe media from the random-write patterns inherent to continuous AI token generation and preserving drive endurance for the lifetime of the cluster.

The Disaggregated Prefill-Decode Architecture

The LLM inference workload is split into two distinct phases with entirely contradictory hardware demands ^[11]:

1. **Prefill Phase:** Ingests the entire prompt simultaneously. This phase is intensely compute-bound, executing massive GEMMs to calculate the initial attention scores.
2. **Decode Phase:** Autoregressively generates output tokens. This phase is severely memory-bandwidth-bound.

Running both phases on the same GPU pool causes catastrophic interference. Consequently, modern deployment pipelines enforce Prefill-Decode (PD) Disaggregation ^[12]. Dedicated prefill nodes process the prompts, generate the KV cache, and physically transmit it across the datacenter network to dedicated decode nodes ^[11]. Systems leverage the Mooncake Transfer Engine and LMCACHE to execute zero-copy GPUDirect RDMA transfers across InfiniBand or RoCE v2 fabrics, achieving up to 142.25 GB/s of peak cross-node transfer bandwidth ^[13].

System Paradigms: Network-Aware Schedulers for Disaggregated Serving

While prefill-decode disaggregation solves compute-memory contention and boosts inference throughput in disaggregated AI factories, it forces the generated KV cache to traverse the datacenter network before decoding begins. This introduces significant network-transfer overhead, directly inflating the Time To First Token (TTFT). Traditional schedulers route requests based solely on compute load and prefix-cache locality, remaining entirely topology-agnostic. To resolve this bottleneck, next-generation schedulers like NetKV target this hurdle by optimizing decode instance selection to minimize network delay. By introducing a "network cost oracle" ^[14], NetKV dynamically evaluates the topological distance and real-time fabric congestion between the prefill publisher and the decode consumer. Implementing a per-request greedy algorithm, NetKV effectively ranks instances by balancing cache locality against network transfer costs. Recent abstract data reveals that NetKV reduces mean TTFT by up to 21.2% over round-robin placements and **17.6% over a tuned cache+load-aware scheduler**, while lifting Service Level Objective (SLO) attainment by **up to 20.1 percentage points** and keeping Time Between Tokens (TBT) overhead strictly below 0.5 ms ^[14].

3. Heterogeneous Compute Vectors, Spatial Architecture Pivots, and Compute-In-Memory (CIM)

Rather than simply climbing a linear 1D hierarchy of capacity versus latency (from host CPU down to storage), next-generation architectures are pivoting into a multi-dimensional matrix. These heterogeneous paradigms fundamentally alter the spatial relationship between the Arithmetic Logic Unit (ALU) and the memory cell, establishing distinct operational vectors:

Vector 1: Ultra-Scaled Tier 0 (SRAM-Dominant Architectures) Standard GPU architectures rely on HBM to act as the primary working memory, but they ultimately hit an off-chip "memory wall" during batch-size-1 inference. To bypass this, accelerators dramatically expand Tier 0 capacity using massive compute-integrated SRAM:

- **Cerebras (WSE-3):** The Wafer-Scale Engine 3 distributes 44 GB of SRAM directly adjacent to 900,000 compute cores across a single massive die, achieving an astounding 21 PB/s of memory bandwidth to circumvent the HBM bottleneck.
- **Groq 3 LPU:** Designed for ultra-low-latency token decode, Groq LP30 chips discard HBM in favor of 500 MB of on-die SRAM per unit, yielding 150 TB/s of deterministic, compiler-driven bandwidth.

Vector 2: Interconnect-Optimized Custom ASICs (TPUs) Google's Tensor Processing Units (TPUs) prioritize native Inter-Chip Interconnects (ICI). By fusing chips together using an 800 GB/s bidirectional 2D torus topology, the TPU v6e (Trillium) architecture treats multi-chip distributed memory as a first-class mathematical construct, prioritizing synchronous pod-level scale-out over isolated single-chip compute density.

Vector 3: Compute-In-Memory (CIM) and Compute-Near-Memory This vector completely reverses the Von Neumann model by integrating Arithmetic Logic Units (ALUs) directly into the memory silicon, effectively eliminating data transit power costs. Standards like JEDEC's LPDDR6-PIM, alongside proprietary designs like Samsung's FIM-DRAM and UPMEM's architecture, embed vector computation native to the DRAM die, allowing matrix operations to execute at the source of the data.

Vector 4: Compute-In-Network and Smart Fabrics Data movement logic is increasingly being offloaded from the host OS directly to the network interface. Utilizing Data Processing Units (DPUs) as "storage brains" (e.g., NVIDIA's BlueField-4), architectures autonomously manage KV cache routing, hardware-accelerated encryption, and Flexible Data Placement across the Ethernet fabric.^[24]

Vector 5: Spatial Pivots and Coherent Agentic Offloading The mechanical layout of memory directly changes spatial system design. By shifting from traditional vertical DIMMs to flat-lying SOCAMM2 modules, server node engineering has executed a foundational spatial pivot. This horizontal footprint drastically improves wind-tunnel aerodynamics and liquid cooling fluid dynamics, increasing per-rack compute density while allowing a single fluid cold-plate to clamp across both the host CPU (e.g., NVIDIA Vera or AMD Verano) and its entire memory subsystem simultaneously. This physical proximity enables high-speed, coherent agentic offloading, allowing the host CPU to leverage its LPDDR5X SOCAMM2 memory pool as a high-speed staging ground to prune and manage the KV cache, isolating the accelerator for raw token generation.

4. Microarchitecture, Hardware-Software Co-Design, and Universal Compilation

To saturate the arithmetic logic capabilities of the GPU, extreme bare-metal hardware-software co-design is required to remove padding overheads and host CPU dependencies.

Universal Compilation and the Abstraction of Heterogeneous Vectors

The proliferation of heterogeneous accelerators—ranging from ultra-scaled SRAM engines (Vector 1) to processing-in-memory arrays (Vector 3)—presents a severe software fragmentation challenge. Traditional, rigid compilation pathways often rely on offline, hardware-specific conversions (such as exporting to ONNX and compiling via proprietary, closed-source toolchains), which fail to dynamically exploit disaggregated fabrics.

To resolve this, the industry is pivoting toward Universal Compilation frameworks and runtimes, such as Modular's MAX engine, MLIR (Multi-Level Intermediate Representation), and OpenAI's Triton. These universal "compile once, run everywhere" models bridge the gap between high-level AI frameworks (like PyTorch) and specialized Domain-Specific Accelerators (DSAs). Crucially, these compilers are tightly co-designed with zero-kernel data pathways. They automatically map semantic program graphs directly to hardware-level PCIe DMA buffers, GPU-initiated I/O queues (like SCADA), and CXL memory spaces, entirely bypassing manual OS-level thread intervention and ensuring that execution logic seamlessly scales across the multi-dimensional taxonomy.

Software-Driven Pipeline Transpositions: FlashMLA-ETAP

NVIDIA GPUs process matrix math using Warpgroup Matrix-Multiply-Accumulate (WGMMMA) instructions, which

strictly require an M-dimension of at least 64 elements. In Multi-Head Latent Attention (MLA) decoding scenarios, the per-GPU head count is frequently reduced due to insufficient memory necessitating head splitting (e.g., splitting the DeepSeek-R1 671B model's 128 heads across eight GPUs, leaving only 16 heads per GPU). This falls below the WGMMA minimum M-dimension, causing massive padding overheads and reducing compute utilization to 25% or lower.

To solve this, frameworks like FlashMLA-ETAP (Efficient Transpose Attention Pipeline) perform software-driven pipeline transpositions. ETAP transposes the layout of the attention mechanism to compute the transposition:

$$\mathbf{S}^T = \mathbf{K} \cdot \mathbf{Q}^T$$

$$\mathbf{P}^T = \text{softmax}(\mathbf{S}^T)$$

$$\mathbf{O} = (\mathbf{V}^T \cdot \mathbf{P}^T)^T$$

This redefines the WGMMA M-dimension as the KV context length (which is naturally large and requires no padding) instead of the query length, while the query-related N-dimension is handled with a flexible 8-element granularity. This co-designed software maneuver eliminates hardware padding overhead entirely, yielding a 2.78x speedup on NVIDIA H20 GPUs. ^[25]

Hardware-Software Co-Design in AI Workloads

The theoretical peak bandwidths and physical topologies of Tier 1 through Tier 4 hardware are ultimately constrained by the software abstractions utilized to address them. In traditional POSIX-compliant environments, data transfer between the Network Interface Card (NIC), the host CPU memory, and persistent storage is paralyzingly inefficient. Standard system calls require expensive context switches from user space into kernel space, triggering hardware interrupts and multiple redundant memory copies. At the scale of terabytes per second, kernel-mediated data movement is a catastrophic bottleneck that induces unacceptable tail latencies. ^[26]

Algorithmic Bandwidth Bypassing: Speculative Decoding and dFlash on Blackwell

The Decode Memory Bottleneck

While structural innovations successfully mitigate prefill computational overheads, the autoregressive mechanics of the LLM decode phase expose an entirely distinct architectural constraint. As Large Language Models expand context windows into the multi-million token regime, the volumetric capacity and bandwidth of the Key-Value (KV) cache have become the primary inhibitors of scalable inference. During the generation cycle, the decode phase is fundamentally memory-bandwidth bound; for every discrete token generated, the GPU memory controllers must retrieve the associated KV cache fragments. This rigid ratio between token production and memory retrieval leaves high-performance arithmetic logic—such as advanced Tensor Cores—critically underutilized, as compute engines are essentially starved by the logistics of data transit.

The Speculative Workaround

To circumvent this pervasive memory wall, advanced hardware-software co-design frameworks employ *Speculative Decoding*. This strategic algorithmic pivot weaponizes abundant, idle compute cycles to insulate and

preserve increasingly scarce memory bandwidth. Rather than executing sequential token generation via the trillion-parameter target model, a compact, high-efficiency "draft" model rapidly predicts a sequence of prospective tokens. The primary target model subsequently ingests this entire draft sequence in a single parallel operation, validating the tokens against the KV cache in one consolidated memory fetch. By aggregating the verification phase, speculative decoding fundamentally minimizes the frequency of high-latency memory round-trips.

Hardware Co-Design: dFlash on Blackwell

The transformative potential of this algorithmic shift is fully extracted through deep integration with low-level microarchitecture, exemplified by the NVIDIA Blackwell architecture utilizing dFlash. By synchronizing speculative draft-and-verify software pipelines with Blackwell's sophisticated tensor core scheduling and memory alignment protocols, dFlash eliminates the latency overhead of draft generation while maximizing parallel verification throughput. This convergence of algorithmic bypassing and specialized silicon achieves a 15x acceleration in inference performance, validating that software-driven pipeline transpositions are as essential as interconnect scaling in dismantling the AI memory wall. ^[27]

GPU-Initiated Storage and Drive-Level Capacity Mathematics

The traditional I/O path forces the host CPU to act as the intermediary for all storage requests. Next-generation systems leverage GPU-Initiated Storage architectures—such as NVIDIA SCADA and Big Accelerator Memory (BaM)—which empower the GPU's memory controllers to issue low-level NVMe commands directly to the SSD pool across the PCIe bus.

However, as local NVMe drives scale into the petabyte range, the on-board DRAM footprint required by the SSD's Flash Translation Layer (FTL) becomes economically unviable. To mitigate this, storage controllers are transitioning from 4K block sizes to 16K or 64K Indirection Units (IUs). This hardware shift drastically shrinks the logical-to-physical (L2P) mapping table, governed by the following memory allocation formula:

$$M_{\text{DRAM}} = \left(\frac{C_{\text{SSD}}}{U_{\text{IU}}} \right) \cdot S_{\text{entry}}$$

Where M_{DRAM} is the required onboard memory, C_{SSD} is the total drive capacity, U_{IU} is the size of the Indirection Unit, and S_{entry} is the byte-size of the mapping entry.

Fixing Synchronous Core Contention

In distributed Graph Neural Network (GNN) and data-intensive AI workloads, legacy GPU-initiated storage systems (such as BaM and GIDS) enabled accelerators to issue NVMe commands directly, bypassing the host CPU ^[28]. However, these early architectures utilized a synchronous I/O stack where a single GPU thread managed both the submission of I/O commands and the polling for their completion. This monolithic approach monopolized massive amounts of GPU compute—often tying up vital thread resources just to manage storage access—which severely constrained parallel computation and limited the ability to overlap arithmetic with data retrieval.

Pipeline-Friendly Asynchronous Disk I/O

To resolve this core contention, the Hyperion architecture completely disaggregates the GPU-initiated I/O process into a pipeline-friendly asynchronous disk I/O stack ^[29, 30]. Hyperion achieves this by splitting the I/O lifecycle into two separate, highly specialized kernels. The first is a thread-level parallel I/O command submission kernel that requires an extremely small fraction of the GPU cores (around ~1%) to fully saturate the NVMe SSD throughput. The second is an asynchronous I/O completion handling kernel that processes fetched data (requiring ~30% of cores). By decoupling these operations, Hyperion frees the vast majority of the GPU's Streaming Multiprocessors (SMs) for actual AI arithmetic, enabling deep CUDA-level pipelining that perfectly overlaps the disk I/O preparation of current data batches with the model training of previous batches.

The TPC Economic Metric

To quantify the enterprise value of this microarchitectural shift, Hyperion introduces the Throughput Per monetary Cost (TPC) metric, defined as the number of training samples processed per second per dollar spent on hardware. Utilizing a TPC-analytical model to guide hardware allocation, benchmarks prove that deploying inexpensive NVMe SSDs on a single node equipped with Hyperion's asynchronous I/O yields massive economic advantages. This localized, NVMe-scaled topology achieves up to a 60x TPC advantage over traditional, highly expensive distributed in-memory clusters like DistDGL, proving that extreme I/O co-design is vastly more cost-efficient than brute-force hardware scaling.

Efficient General Intelligence (EGI) and Hierarchical Memory Transformers

The Biological Imperative

The brute-force scaling of Large Language Models (LLMs) has triggered an unsustainable power trajectory, fundamentally motivating the shift toward Efficient General Intelligence (EGI) ^[31]. Serving a massive model like DeepSeek-R1 across an array of 32 NVIDIA H200 GPUs consumes approximately 11.2 kW of power. Over a 12-hour period, this infrastructure burns 483 megajoules (MJ) of energy—equivalent to the power consumed by an average US household over four entire days. In stark contrast, the human brain operates at a highly efficient 12-20W. A human executing 12 hours of rigorous research work consumes only 518 kilojoules (kJ) of energy, which is the caloric equivalent of just two bananas. This biological imperative dictates that future AI architectures must abandon flat, energy-intensive memory scaling in favor of extreme algorithmic efficiency, necessitating deep model and hardware co-optimization ^[31].

Hierarchical Memory Transformers (HMT)

To mimic this biological efficiency, Dr. Jason Cong's research at UCLA introduces the Hierarchical Memory Transformer (HMT). Inspired by the human brain's ability to selectively recall past experiences without retaining every granular detail, HMT breaks away from flat context windows by stratifying memory into distinct tiers: a sensory/active context tier, a short-term tier utilizing the GPU's L2 Cache and HBM, and a long-term retention tier anchored by dense SSDs.

Algorithmic Compression

Standard transformer architectures suffer a massive computational penalty because they must execute $O(n^2)$ attention over 1M+ input tokens for every pass. HMT avoids this brute-force processing through aggressive algorithmic compression and relevance checking. Instead of retaining every token, the system compresses previous segments into highly dense memory embeddings (e.g., 4096 FP32 values per summary). These embeddings capture both the current context and past information, reducing information loss while allowing the model to retrieve only the relevant past segments via cross-attention lookups.

Hardware Impact

This algorithmic pivot radically alters the hardware demands of the memory wall. Benchmarks comparing HMT-optimized models against standard large long-context models show a 7x to 20.9x reduction in memory footprint requirements. For example, while a traditional ~20B parameter model requires 46.7 GB of VRAM and consumes 137W of power to process long sequences, an HMT-augmented architecture can process equivalent summarization and question-answering tasks using as little as 6.1 GB of memory. In the most optimized configurations, HMT slashes the footprint down to just 0.75 GB (a 20.9x reduction), yielding massive subsequent reductions in power consumption (dropping from 120W to just 35W) and proving that algorithmic memory compression is an essential prerequisite for reaching EGI.

Democratizing Chip Design with AI and Automated Transformation

This paradigm shift allows software engineers to build domain-specific accelerators (DSAs) without extensive VLSI or EDA expertise. Frameworks pioneered by experts like Professor Jason Cong (UCLA VAST Lab) ^[15] (UCLA VAST Lab) use AI to automate optimization, turning specialized physical layout workflows into automated code compilation pipelines.

The Core Challenge

- **Talent Gap** ^[16]: Hardware designers are significantly outnumbered by software developers globally.
- **Complexity**: Custom silicon design typically requires deep hardware knowledge.
- **High Costs** ^[17]: Traditional EDA tools involve massive investments and multi-year timelines.

Key Technical Pillars

The transition to an accessible ecosystem relies on High-Level Synthesis (HLS) and AI algorithms:

- **High Level Synthesis (HLS)** ^[18]: Translates C/C++ code directly into hardware description layouts (RTL).
- **Automated Transformation** ^[18]: ML compilers automate directive placement and loop optimization, replacing human intuition.
- **Graph Neural Networks (GNNs)** ^[19]: Analyze circuit structures early to predict performance, power, and area (PPA).
- **Design Space Exploration (DSE)** ^[19]: RL agents automatically explore millions of hardware combinations to find optimal designs.

Industry Impact

3. **Faster Innovation** ^[20]: Prototyping timelines drop from months to weeks.
4. **Accessibility** ^[21]: Startups and researchers can manufacture specialized silicon economically.
5. **Co Design** ^[22, 23]: Silicon layouts can adapt fluidly to changing software and neural network structures.

5. Neuromorphic Systems, Infrastructure Telemetry, and the RAS Gap Challenge

The physical limits of scaling traditional Von Neumann architectures have driven a fundamental paradigm shift as systems approach the scaling limits of Artificial General Intelligence (AGI).

Non-Von Neumann Paradigms and the AGI Event Horizon

The transition toward autonomous, multi-turn reasoning is shifting the industry from static "Cerebral AI" to dynamic

"Embodied AI" systems. This bio-digital convergence demands the absolute abandonment of sequential Von Neumann hardware in favor of asynchronous, neuromorphic substrates. Modeled after the hippocampal-entorhinal structures of the biological brain, these architectures co-locate memory and logic within networks of artificial neurons and synapses that communicate via sparse, electrical "spikes." Digital neuromorphic accelerators (such as Intel's Loihi 2) resolve physical routing challenges via MLIR-AIE compilation paths that compile semantic program graphs directly to physical tile DMA buffers. Combined with edge accelerators like the 28nm Spiking Vision Transformer (Spiking ViT), these non-Von Neumann structures achieve millisecond-level, zero-copy sense-act loops.

Infrastructure Telemetry and Power Management

To support these ultra-dense computing environments, data center power management requires radical transparency. Modern deployments operating PCIe Gen6 and Gen7 liquid-cooled racks utilize advanced NVMe telemetry—specifically the **TP4199 Device Self Power Monitoring** and **TP4210 Device Voltage Monitor** proposals^[51, 52]. These enhancements provide granular visibility into SSD power behavior and host supply voltages, enabling active thermal management and allowing infrastructure engineers to transition from reactive troubleshooting to proactive failure prediction without exposing end-user data.

The RAS Gap Challenge

Despite these advancements, infrastructure architects must navigate a critical hurdle: the adoption of mobile-derived LPDDR memory in enterprise data centers. Because LPDDR was historically designed for consumer devices, scaling it to handle mission-critical AI data pipelines introduces the "RAS Gap Challenge." To meet enterprise Reliability, Availability, and Serviceability (RAS) standards, LPDDR6 modules require ongoing hardware enhancements, including Per Row Activation Counting (PRAC) to combat Rowhammer attacks, Carve-out Meta modes for data isolation, programmable link protection schemes, and robust on-die multi-bit error correction code (ECC)^[53].

6. Emerging Topologies Deep Dive: LPDDR6, HBM4, and Beyond

High Bandwidth Memory 4 (HBM4)

Ratified in April 2025 by the JEDEC Solid State Technology Association as standard JESD270-4, the HBM4 specification represents a dramatic shift from its predecessors by doubling the physical interface width.^[32] Where HBM3e utilized a 1024-bit interface, HBM4 scales to an unprecedented 2048-bit interface per stack, divided into 32 independent channels (and 64 pseudo-channels).^[32] Operating at per-pin data rates ranging from 6.4 Gbps up to 13 Gbps, a single HBM4 stack delivers extraordinary theoretical bandwidth, calculated as:

$$\text{Bandwidth}_{\text{HBM4}} = \frac{13 \text{ Gbps} \times 2048 \text{ bits}}{8 \text{ bits/byte}} = 3.328 \text{ TB/s}$$

HBM4 eliminates traditional microbump packaging in favor of direct logic-to-memory hybrid bonding, enabling tighter integration with compute dies and reducing latency.^[33] The standard allows stack heights to scale from 4-high up to 16-high, yielding capacities up to 64 GB per individual package.^[32] Core voltage is lowered to a VDDQ of 0.7V to 0.9V and a VDDC of 1.05V, pulling the energy efficiency down to the highly desirable ~1.5 pJ/bit threshold

(representing a 60% power efficiency gain over HBM2/2E).^[34]

Low-Power Double Data Rate 6 (LPDDR6) and SOCAMM2

LPDDR6 completely redesigns the standard channel architecture, standardizing on a 24-bit channel constructed from two highly granular 12-bit sub-channels with 12 data signal lines (DQs) per sub-channel.^[35] Operating at introductory speeds of 10.7 Gbps and scaling to defined maximums of 14.4 Gbps, a single LPDDR6 channel produces massive throughput:

$$\text{Bandwidth}_{\text{LPDDR6}} = \frac{14.4 \text{ Gbps} \times 24 \text{ bits}}{8 \text{ bits/byte}} = 43.2 \text{ GB/s peak (raw)} \approx 38.4 \text{ GB/s effective}$$

While HBM4 handles the GPU-bound execution of multi-petaflop matrix operations, the host CPU orchestration layer requires a vastly different memory profile. Low-power mobile memory is transitioning into the server space via the JEDEC-standardized open SOCAMM2 (System On Chip Advanced Memory Module 2) ecosystem.^[36] NVIDIA has aggressively driven the proprietary validation of this standard for its Vera CPU. Designed with 88 custom Olympus cores, the Vera architecture pairs with up to 1.5 TB of LPDDR5X SOCAMM2, pushing a massive 1.2 TB/s of aggregate memory bandwidth. AMD's 2027 roadmap confirms that select 6th Gen EPYC "Verano" SKUs will similarly deploy LPDDR5X SOCAMM2 to optimize performance-per-system-watt as a host CPU for future Instinct GPU rack-scale solutions.^[37]

Interconnect Wars: UALink vs. Proprietary Fabrics

As memory pooling scales across racks, the proprietary nature of fabrics like NVIDIA's NVLink and NVSwitch has prompted an open-standard response: the Ultra Accelerator Link (UALink). Finalized in April 2025, the UALink 1.0 specification defines an open memory-semantic interconnect built on Ethernet physicals, designed specifically for AI computing pods. Operating at 200G per lane, UALink supports native read, write, and atomic memory operations directly between accelerators and switches, enabling a scale-up connection for up to 1,024 accelerators within a single shared memory domain.

In parallel, the CXL Consortium released CXL 4.0 in November 2025, doubling bandwidth to 128 GT/s by integrating the PCIe 7.0 physical layer. CXL 4.0 introduces "Bundled Ports"—an architectural enhancement that aggregates multiple upstream physical CXL device ports into a single, logical entity, enabling link bandwidths of up to 1.5 TB/s.

Interconnect Latency Analysis: NVLink/NVSwitch vs. PCIe

To understand why advanced AI clusters aggressively move away from commodity standard buses for distributed operations, it is critical to perform a microarchitectural latency comparison. The following matrix compares the physical propagation delays, protocol processing steps, and software overheads that define the performance of proprietary fabrics versus PCIe configurations:

System Operation / Interconnect Path	Typical Latency (ns)	Typical Latency (μs)	Architectural Context & Latency Contributors
HBM4 Stack Access	10 - 30 ns	0.01 - 0.03 μs	Direct on-package silicon access via high-density TSVs

PCIe Gen5 Link FEC (Per Direction)	2 - 4 ns	0.002 - 0.004 μ s	Hardware forward error correction penalty
PCIe Gen6 Link FEC (Per Direction)	4 - 8 ns	0.004 - 0.008 μ s	Stronger correction required to maintain PAM4 signal integrity
PCIe Gen5 Switch Port-to-Port	115 ns	0.115 μ s	Low-latency P2P port-to-port bypass (e.g., Broadcom PEX89048)
NVLink 4.0/5.0 Direct (Intra-Node P2P)	100 - 200 ns	0.1 - 0.2 μ s	Point-to-point GPU memory-semantic hardware links
NVSwitch Traversal (Single Hop)	150 - 300 ns	0.15 - 0.3 μ s	Central non-blocking crossbar routing inside DGX node
Direct GPU-to-GPU Memory Copy (NVLink)	2,000 ns	2 μ s	Unified Memory/NCCL-optimized peer-to-peer data copy
PCIe Gen5 GPU-to-GPU (Direct P2P)	1,000 - 2,000 ns	1 - 2 μ s	P2P transfer traversing CPU PCIe root complex
NVSwitch Traversal (2-Hop Cross-Pod)	4,000 - 5,000 ns	4 - 5 μ s	Scaling out to 72+ GPU racks using multi-hop NVSwitch fabrics
GPU-to-GPU Memory Copy (PCIe Gen4)	20,000 ns	20 μ s	Standard transfer blocked by operating system driver copy loops

This rule-of-thumb analysis demonstrates that NVLink provides an approximate 10x reduction in peer-to-peer memory copy latency compared to standard PCIe channels. This latency reduction is primarily driven by three factors: the avoidance of intermediate host CPU memory hops, the elimination of multi-microsecond software-layer copy latencies via hardware-supported Address Translation Services (ATS), and the optimization of on-package die-to-die interfaces (like NVLink-C2C) that eliminate board-level propagation delays entirely.

7. Synthesis and Conclusion

The empirical evidence derived from current-generation hardware specifications, emerging interconnect standards, and cutting-edge software frameworks confirms a foundational paradigm shift in systems architecture. The transition toward agentic AI and trillion-parameter models has decisively exposed the limitations of traditional Von Neumann scaling, rendering the strict, isolated boundaries of legacy memory and storage hierarchies obsolete. To navigate this complexity, this whitepaper establishes a comprehensive blueprint for the modern AI factory, structured around the following core architectural pillars:

- The Disaggregated Storage & Memory Fabric:** The historical compute-storage boundary has dissolved into a continuous, linear tiering system. The formal integration of CXL 3.1/4.0 fabrics unifies these tiers

semantically via CXL.mem. Crucially, the introduction of Tier 2.5 (G2.5) Persistent Context Memory, powered by CXL Type-3 persistent memory devices, establishes a vital, byte-addressable staging layer for **flash-wear-free** (bypassing NAND flash wear by utilizing the CXL memory tier) transaction logging and KV cache persistence, bridging the gap between volatile DRAM and block-addressable NVMe storage.

- **Software-Driven Pipeline Transpositions:** To maximize hardware utilization, frameworks are executing fundamental pipeline transpositions. Techniques like FlashMLA-ETAP eliminate hardware padding overheads by aligning memory layouts with WGMMMA requirements, while Speculative Decoding (integrated with Blackwell-era dFlash) weaponizes compute cycles to maximize memory bandwidth efficiency, collectively dismantling the inference memory wall.
- **Zero-Kernel Data Planes:** Sustaining millions of IOPS requires the total elimination of the OS kernel. Modern architectures leverage user-space polling (SPDK), kernel-bypass networking (DPDK), and GPU-initiated I/O architectures (like Hyperion and BaM) to ensure end-to-end zero-copy routing, allowing for asynchronous, pipeline-friendly I/O that overlaps arithmetic with data retrieval.
- **Heterogeneous Architectural Vectors:** Systems are pivoting into a multi-dimensional matrix. SRAM-dominant architectures (e.g., Groq LPU, Cerebras) bypass HBM bottlenecks, while Compute-in-Memory (CIM) vectors (LPDDR6-PIM) decentralize operations to the memory die. Furthermore, horizontal spatial pivots like SOCAMM2 improve thermal dynamics, enabling host CPUs (like the Vera CPU) to execute coherent agentic offloading via NVLink-C2C.
- **Neuromorphic Frontiers & RAS Management:** As AI scales, architectures are shifting toward asynchronous neuromorphic substrates for embodied reasoning. This necessitates robust RAS management for enterprise LPDDR6 deployment, ensuring reliability through PRAC (Per Row Activation Counting) and hardened error correction.

Ultimately, demolishing the "Memory Wall" demands that data center infrastructure be treated as a single, vertically integrated machine. By embracing CXL-scale-out topologies, bare-metal data planes, and heterogeneous co-design, systems engineers can ensure that structural performance scales seamlessly with the cognitive demands of the agentic AI era.

Reference:

1. AI and Memory Wall, <https://arxiv.org/html/2403.14123v1>
2. GPU Memory Essentials for AI Performance, <https://developer.nvidia.com/blog/gpu-memory-essentials-for-ai-performance/>
3. Famfs: open source scale-out shared memory file system - MSST, accessed June 20, 2026, [https://msstconference.org/MSST-history/2025/Presentations/Day%20%20\(September%2023rd\)/Session%201%20\(pNFS\)/famfs-Micron-msst-final.pdf](https://msstconference.org/MSST-history/2025/Presentations/Day%20%20(September%2023rd)/Session%201%20(pNFS)/famfs-Micron-msst-final.pdf)
4. Introducing the CXL 3.1 Specification - Compute Express Link, https://computeexpresslink.org/wp-content/uploads/2024/03/CXL_3.1-Webinar-Presentation_Feb_2024.pdf
5. System Composability Using CXL, <https://www.openfabrics.org/wp-content/uploads/2024-workshop/2024-workshop-presentations/session-10.pdf>
6. Eliminating the OS Kernel from the Data Plane, A Survey of Architectural Paradigms for User-Space Storage Engines and Implementation Strategy for Dataplane-Emu, <https://github.com/SiliconLanguage/dataplane-emu/blob/main/docs/publications/eliminating-os-kernel/eliminating-os-kernel.pdf>
7. Understanding host network stack overheads - ResearchGate, https://www.researchgate.net/publication/353785836_Understanding_host_network_stack_overheads
8. Next-Level Inference: Why Your Single-Node vLLM Setup Needs Prefill-Decode Disaggregation, <https://vllm.ai/blog/2026-04-07-moriio-kv-connector>
9. Understanding CMX (Context Memory eXtension) in Data Storage for AI - Solidigm, accessed June 20, 2026, <https://www.solidigm.com/products/technology/what-is-cmx-context-memory-storage.html>
10. NVIDIA CMX Context Memory Storage Platform, <https://www.nvidia.com/en-us/data-center/ai-storage/cmx/>
11. Proactive Intra-GPU Disaggregation of Prefill and Decode in LLM Serving, <https://arxiv.org/pdf/2507.06608v4>

12. Designing distributed AI inference: Core concepts and scaling dimensions, <https://developers.redhat.com/articles/2026/06/22/designing-distributed-ai-inference-core-concepts-and-scaling-dimensions>
13. Combating the Memory Walls: Optimization Pathways for Long-Context Agentic LLM Inference, <https://arxiv.org/pdf/2509.09505>
14. A Strategic Briefing on Memory Tiering, CXL Disaggregation, and Agentic Compute Capitalization, <https://github.com/SiliconLanguage/tensorplane/blob/main/docs/whitepapers/demolishing-the-memory-wall/Demolishing%20the%20Memory%20Wall%20for%20AGI-Scale%20LLM%20Inference.pdf>
15. NetKV: Network-Aware Decode Instance Selection for Disaggregated LLM Inference, <https://arxiv.org/pdf/2606.03910>
16. LMCACHE: An Efficient KV Cache Layer for Enterprise-Scale LLM Inference - arXiv, accessed June 20, 2026, <https://arxiv.org/html/2510.09665v2>
17. SiliconLanguage - GitHub, accessed June 20, 2026, <https://github.com/SiliconLanguage>
18. FlashMLA-ETAP: Efficient Transpose Attention Pipeline for Accelerating MLA Inference on NVIDIA H20 GPUs - arXiv, accessed June 20, 2026, <https://arxiv.org/html/2506.01969v3>
19. From data to intelligence: Micron's role in the AI revolution, <https://www.micron.com/about/blog/applications/ai/from-data-to-intelligence-microns-role-in-the-ai-revolution>
20. Boost Inference Performance up to 15x on NVIDIA Blackwell Using dFlash Speculative Decoding. NVIDIA Technical Blog, accessed June 20, 2026, <https://developer.nvidia.com/blog/boost-inference-performance-up-to-15x-on-nvidia-blackwell-using-dflash-speculative-decoding/>
21. What Is Context Memory? The New Infrastructure Layer Powering AI Inference - WEKA, accessed June 20, 2026, <https://www.weka.io/learn/ai-ml/what-is-context-memory-primary/>
22. A Survey on User-Space Storage and Its Implementations, <https://arxiv.org/html/2306.10503v1>
23. KV CACHE OPTIMIZATION STRATEGIES FOR SCALABLE AND EFFICIENT LLM INFERENCE, <https://arxiv.org/pdf/2603.20397v1>
24. fabric-lib: RDMA POINT-TO-POINT COMMUNICATION FOR LLM SYSTEMS, <https://arxiv.org/pdf/2510.27656>
25. NetKV: Network-Aware Decode Instance Selection for Disaggregated LLM Inference, <https://arxiv.org/pdf/2606.03910>
26. AI Is About to Democratize Chip Design — Cognichip CEO Interview, <https://www.youtube.com/watch?v=o8nFanOouw0&vl=en>
27. Workshops on Extreme Scale Design Automation (ESDA) Challenges and Opportunities for 2025 and Beyond, <https://arxiv.org/pdf/2005.01588>
28. Report for NSF Workshop on AI for Electronic Design Automation, <https://arxiv.org/html/2601.14541v4>
29. AI for Electronic Design Automation (EDA) Workshop, <https://ai4eda-workshop.github.io/>
30. Faster Innovation, <https://www.eejournal.com/article/the-eda-enigma/>
31. Halfacree, G. (2022). Google, SkyWater, Efabless Announce a New Open-Silicon PDK on a 90nm Process Node. Hackster.io. Retrieved from <https://www.hackster.io/news/google-skywater-efabless-announce-a-new-open-silicon-pdk-on-a-90nm-process-node-c3e644a9529d>
32. The Democratization of Co-design by James Ang and Antonino Tumeo, <https://www.youtube.com/watch?v=h0dHD3mEYal>
33. ChipGPT: The AI Breakthrough That's Redefining Semiconductor Design, https://www.linkedin.com/posts/kishorakshinthala_hdls-vhdl-llm-activity-7295490610327687168-1arr
34. JEDEC Officially Releases HBM4 Memory Standard - News - All About Circuits, accessed June 20, 2026, <https://www.allaboutcircuits.com/news/jedec-officially-releases-hbm4-memory-standard/>
35. High Bandwidth Memory - Grokipedia, accessed June 20, 2026, https://grokipedia.com/page/High_Bandwidth_Memory
36. HBM3e and HBM4: IC design guide for next-generation high bandwidth memory, accessed June 20, 2026, <https://blogs.sw.siemens.com/semiconductor-packaging/2026/04/24/hbm3e-hbm4-ic-design-guide/>
37. Business Wire: "JEDEC Releases New LPDDR6 Standard to Enhance Mobile and AI Memory Performance" : r/hardware - Reddit, accessed June 20, 2026, https://www.reddit.com/r/hardware/comments/1lvpmww/business_wire_jedec_releases_new_lpddr6_standard/
38. [News] JEDEC Previews LPDDR6 Enhancements, Develops SOCAMM2 Standard for AI Memory - TrendForce, accessed June 20, 2026, <https://www.trendforce.com/news/2026/04/24/news-jedec-previews-lpddr6-enhancements-develops-socamm2>

[-standard-for-ai-memory/](#)

39. JEDEC previews LPDDR6, confirming its datacenter focus - OC3D, accessed June 20, 2026, <https://overclock3d.net/news/memory/jedec-previews-lpddr6-proving-that-datacenters-have-stolen-the-mobile-memory-standard/>
40. GPU-Initiated On-Demand High-Throughput Storage Access in the BaM System Architecture, <https://arxiv.org/pdf/2203.04910>
41. Hyperion: Optimizing SSD Access is All You Need to Enable Cost-efficient Out-of-core GNN Training, <https://wangzeke.github.io/doc/Hyperion-ICDE25.pdf>
42. Efficient General Intelligence with Novel Model and Customized Silicon Co-Design, Jason Cong (UCLA), <https://www.youtube.com/watch?v=uAmfjizj69M>
43. Special Joint Engineering and AI Seminar: Jason Cong, UCLA, <https://www.youtube.com/watch?v=Nwkb2NASJpA>
44. Vera Rubin Decoded Pt. 2 | Rubin GPU - Engineering Deep Dive, <https://transurfing-volatility.ghost.io/vera-rubin-decoded-pt2/>
45. Vera Rubin Decoded Pt. 3 | Vera CPU & the Networking Silicon Family, <https://transurfing-volatility.ghost.io/vera-rubin-decoded-pt3/>
46. Vera Rubin Decoded Pt. 5 | Rack Power and the Networking Fabric, <https://transurfing-volatility.ghost.io/vera-rubin-decoded-pt5/>
47. Vera Rubin Decoded Pt. 6 | Supply Chain - Master Reference, <https://transurfing-volatility.ghost.io/vera-rubin-decoded-pt6/>
48. PCI Express® Base Specification Revision 6.4, PCI Special Interest Group, <https://pcisig.com/PCIExpress/Spec/Base/6.4>
49. Low-overhead General-purpose Near-Data Processing in CXL Memory Expanders, IEEE MICRO, <https://ieeexplore.ieee.org/document/10764494>
50. UALink, Ultra Ethernet, and PCIe: Transforming Next-Generation HPC and AI Workloads, Journal of Information Systems Engineering and Management, <https://jisem-journal.com/index.php/journal/article/view/14455/6952>
51. NVM Express Technical Proposal 4199 (TP4199): Device Self Power Monitoring (DSPM) / Self-Reported Drive Power Support, NVM Express, Inc. <https://nvmexpress.org/wp-content/uploads/NVM-Express-Revision-Changes-2025.08.01.pdf>
52. NVM Express Technical Proposal 4210 (TP4210): Advanced Telemetry and Infrastructure Power Management, NVM Express, Inc. (Draft/Ratified Proposal)
53. LPDDR6: The Next-Generation LPDDR Device Standard and How It Differs from LPDDR5 https://community.cadence.com/cadence_blogs_8/b/fv/posts/lpddr6-next-generation-lpddr-device-standard-and-how-it-differs-from-lpddr5
54. CXL Consortium, Introducing the CXL 3.X Specification, https://computeexpresslink.org/wp-content/uploads/2025/02/CXL_Q1-2025-Webinar-Presentation_FINAL.pdf
55. CXL Consortium, Introducing the CXL 4.0 Specification, https://computeexpresslink.org/wp-content/uploads/2025/12/CXL_4.0-Webinar_December-2025_FINAL.pdf
56. Persistent Memory: A Survey of Programming Support and Implementations, <https://dl.acm.org/doi/epdf/10.1145/3465402>
57. Intel Corp., Programming Persistent Memory with PMDK: A Guide for Developers, [PMDK Documentation](#).
58. NVM Express, Flexible Data Placement (FDP) Specification: Optimizing NAND Utilization in Hyperscale AI Factories, [NVM Express Technical Brief](#).
59. SC25 performance breakthrough: 230M IOPS in a single server, <https://www.micron.com/about/blog/storage/ssd/sc25-performance-breakthrough-230m-iops-in-a-single-server>